

# **Neural Networks learn Representation Theory: Reverse Engineering how Networks perform Group Operations**

Bilal Chughtai, Lawrence Chan, Neel Nanda



# PROGRESS MEASURES FOR GROKING VIA MECHANISTIC INTERPRETABILITY

**Neel Nanda**

Independent

neelnanda27@gmail.com

**Lawrence Chan**

UC Berkeley

chanlaw@berkeley.edu

**Tom Lieberum**

Independent

tlieberum3141@gmail.com

**Jess Smith**

Independent



smith.jessk@gmail.com

**Jacob Steinhardt**

UC Berkeley

jsteinhardt@berkeley.edu

# Mystery: Why do models grok?



# GROKking: GENERALIZATION BEYOND OVERFITTING ON SMALL ALGORITHMIC DATASETS


Alethea Power, Yuri Burda, Harri Edwards, Igor Babuschkin  
OpenAI

Vedant Misra\*  
Google

**Methodology:** Apply mechanistic interpretability

# Inspiration: Mechanistic Interpretability

- **Goal:** Reverse engineer neural networks
- **Hypothesis:** Models learn human-comprehensible algorithms and can be understood, if we learn how to make it legible
- Models learn **circuits**, algorithms encoded in the weights
- A deep knowledge of circuits is crucial to understand and predict model behaviour



# Universality

Curve detectors

ALEXNET

Krizhevsky et al. [34]




INCEPTIONV1

Szegedy et al. [26]



VGG19

Simonyan et al. [35]






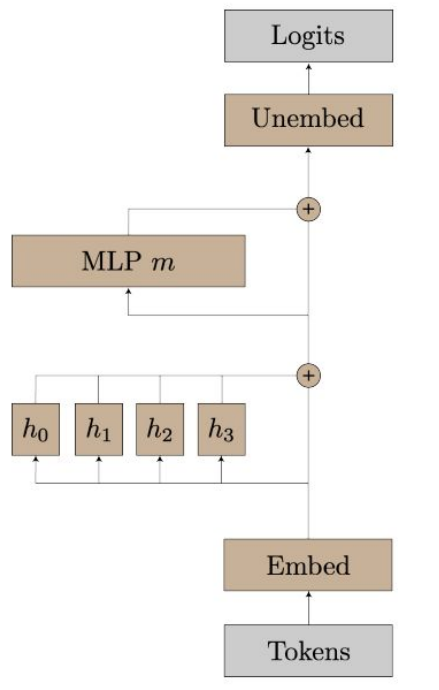
RESNETV2-50

He et al. [36]




High-Low Frequency detectors






Computes logits using further trig identities:

$$\begin{aligned}\text{Logit}(c) &\propto \cos(w(a + b - c)) \\ &= \cos(w(a + b)) \cos(wc) + \sin(w(a + b)) \sin(wc)\end{aligned}$$




Calculates sine and cosine of  $a + b$  using trig identities:

$$\begin{aligned}\sin(w(a + b)) &= \sin(wa) \cos(wb) + \cos(wa) \sin(wb) \\ \cos(w(a + b)) &= \cos(wa) \cos(wb) - \sin(wa) \sin(wb)\end{aligned}$$



Translates one-hot  $a, b$  to Fourier basis:

$$\begin{aligned}a &\rightarrow \sin(wa), \cos(wa) \\ b &\rightarrow \sin(wb), \cos(wb)\end{aligned}$$





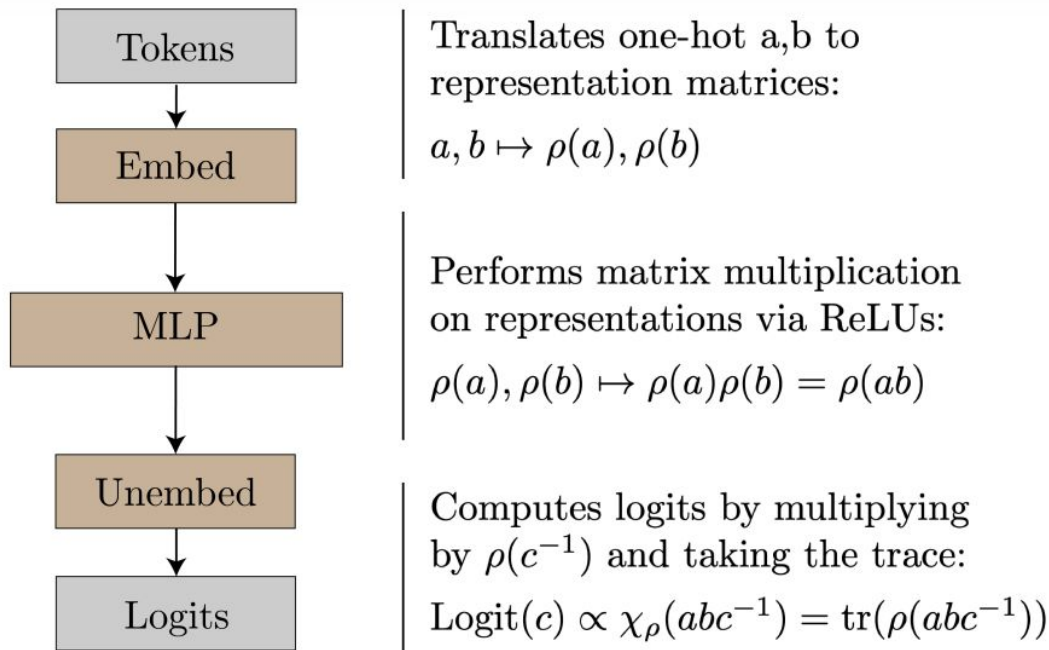
# Representation Theory

A (real) **representation** is a homomorphism, i.e. a map preserving the group structure,  $\rho : G \rightarrow GL(\mathbb{R}^d)$  from the group  $G$  to a  $d$ -dimensional general linear group, the set of invertible square matrices of dimension  $d$ . Representations are in general *reducible*, in a manner we make precise in the Appendix. For each group  $G$ , there exist a finite set of fundamental **irreducible representations**. The **character** of a representation is the trace of the representation  $\chi_\rho : G \rightarrow \mathbb{R}$  given by  $\chi_\rho(g) = \text{tr}(\rho(g))$ . A key fact our algorithm depends on is that character's are maximal when  $\rho(g) = I$ , the identity matrix (Theorem C.8). In particular, the character of the identity element,  $\chi_\rho(e)$ , is maximal.

**Example.** The cyclic group  $C_n$  is generated by a single element  $r$  and naturally represents the set of rotational symmetries of an  $n$ -gon, where  $r$  corresponds to rotation by  $2\pi/n$ . This motivates a 2 dimensional representation – a set of  $n$   $2 \times 2$  matrices, one for each group element:

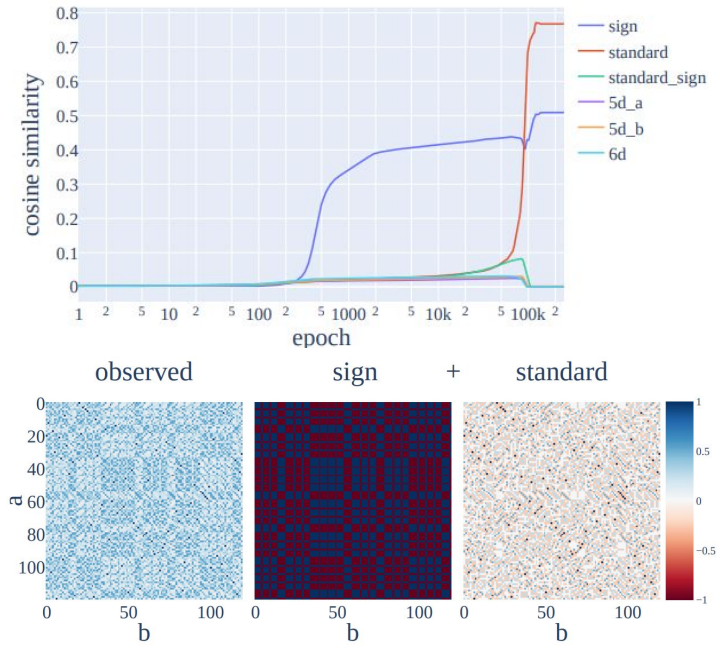
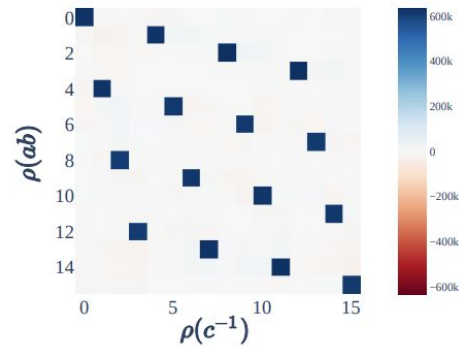
$$\rho(r^k) = \begin{pmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{pmatrix}$$

for element  $r^k$  corresponding to rotation by  $\theta = 2\pi k/n$ .



# Reverse Engineering S5

- 1. Logit similarity
- 2. Embeddings
- 3. MLP activations & the MLP - Logit map
- 4. Ablations



	$W_a$	$W_b$	$W_U$
SIGN	6.95%	6.95%	9.58%
STANDARD	93.0%	93.0%	84.5%
RESIDUAL	0.00%	0.00%	5.96%

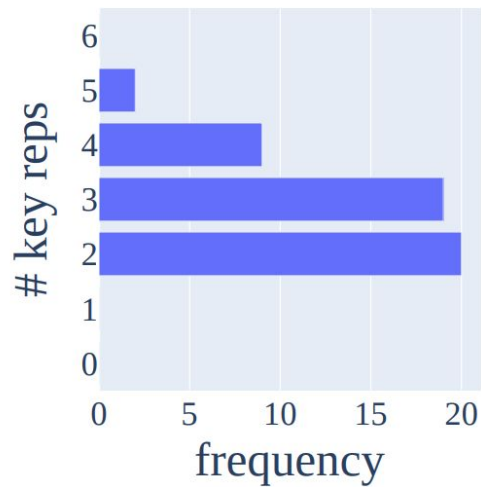
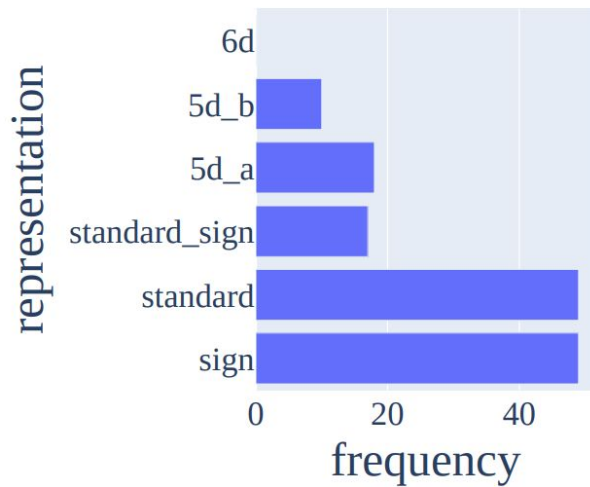
CLUSTER	$\rho(a)$	$\rho(b)$	$\rho(ab)$	RESIDUAL
SIGN	33.3%	33.3%	33.3%	0.00%
STANDARD	39.6%	37.1%	11.3%	12.1%

# Weak Universality

Table 3. Results from all groups on both MLP and Transformer architectures, averaged over 4 seeds. We find that that features for matrices in the key representations are learned consistently, and explain almost all of the variance of embeddings and unembeddings. We find that terms corresponding to  $\rho(ab)$  are consistently present in the MLP neurons, as expected by our algorithm. Excluding and restricting to these terms in the key representations damages performance/does not affect performance respectively.

Group	MLP								Transformer						
	FVE					Loss			FVE				Loss		
	$W_a$	$W_b$	$W_U$	MLP	$\rho(ab)$	Test	Exc.	Res.	$W_E$	$W_L$	MLP	$\rho(ab)$	Test	Exc.	Res.
$C_{113}$	99.53%	99.39%	98.05%	90.25%	12.03%	1.63e-05	5.95	6.88e-03	95.18%	99.52%	92.12%	16.77%	2.67e-07	9.42	2.12e-02
$C_{118}$	99.75%	99.74%	98.43%	95.84%	13.26%	5.39e-06	8.72	3.60e-03	94.05%	99.64%	94.63%	17.11%	1.73e-07	15.93	2.55e-01
$D_{59}$	99.71%	99.73%	98.52%	87.68%	12.44%	6.34e-06	12.37	1.60e-06	98.58%	98.53%	85.01%	10.85%	3.20e-06	46.42	2.82e-05
$D_{61}$	99.26%	99.45%	98.26%	87.61%	12.48%	1.79e-05	12.00	1.69e-06	98.33%	97.40%	85.59%	11.11%	1.63e-02	41.64	9.60e-02
$S_5$	100.00%	99.99%	94.14%	88.91%	12.13%	1.02e-05	11.72	2.21e-07	99.84%	99.97%	85.28%	10.23%	1.43e-07	17.77	4.44e-09
$S_6$	99.65%	99.78%	93.67%	86.38%	8.98%	4.95e-05	12.17	2.66e-06	99.94%	99.93%	86.32%	9.35%	2.21e-06	291.67	1.05e-06
$A_5$	99.04%	99.31%	93.27%	86.69%	10.26%	1.94e-05	9.82	5.28e-07	97.53%	97.40%	83.56%	8.22%	4.88e-02	19.76	7.70e-04

# Strong Universality



# Implications

- Reverse engineering a single network is insufficient for understanding behaviour in general
- BUT it may be possible to build a periodic table of 'universal' features, that in aggregate may be able to explain a given behaviour fully.

## Further Work

- Reverse engineering more group theoretic tasks
- Understanding universality better in algorithmic / realistic tasks
- Understanding network inductive biases better

# Key Takeaways

- Models naturally learn representation theory
- Only by employing the tools of mechanistic interpretability were we able to figure this out

# Future Work

- Further group theoretic tasks
  - Preliminary work has suggested models learn representations in a wider family of group theoretic tasks than simply group composition
- Understanding inductive biases of neural networks better
- Further investigation of universality in algorithmic tasks
- Investigation of universality in realistic tasks and models